

# **Project Report: *AutoJudge – Programming Problem Difficulty Predictor***

## **1. Introduction & Problem Statement**

With the rapid growth of online programming platforms, automatic evaluation of problem difficulty has become an important task. Difficulty labels such as Easy, Medium, and Hard help learners choose appropriate problems and assist platforms in problem curation.

Traditionally, difficulty is assigned manually by problem setters, which can be subjective and time-consuming. This project aims to build an automated system that predicts the difficulty of a programming problem purely from its textual description.

We propose AutoJudge, an NLP-based system that:

1. Predicts a continuous difficulty score (regression)
2. Predicts a difficulty class (Easy / Medium / Hard) (classification)

A web interface is provided using Streamlit to demonstrate real-time predictions.

## **2. Dataset Description**

The dataset used in this project was provided as part of the project assignment.

Each problem contains the following fields:

- Problem title
- Problem description
- Input description
- Output description

- Difficulty score (numeric)
- Difficulty class (Easy / Medium / Hard)

## Dataset Characteristics

- Text-heavy data with algorithmic keywords
- Class imbalance between Easy, Medium, and Hard problems
- Numeric difficulty scores used for regression

	proportion
problem_class	
hard	0.472661
medium	0.341434
easy	0.185905
dtype: float64	

The dataset was split into training and testing sets for both classification and regression tasks.

## 3. Data Preprocessing

All text fields were first **combined into a single text block** per problem to capture full semantic information.

### Preprocessing steps:

- Lowercasing text
- Missing Value Check
- Removing unnecessary whitespace
- Preserving mathematical and algorithmic symbols
- Combining title, description, input, and output sections
- Converted all text to consistent encoding

No aggressive text cleaning was applied to avoid losing algorithm-specific information.

## 4. Feature Engineering

A total of **515 features** were used, consisting of **500 TF-IDF features** and **15 handcrafted features** capturing structural and algorithmic complexity.

## 4.1 TF-IDF Features

- TF-IDF vectorization applied on combined problem text
- Captures important words such as:
  - *dp, graph, constraint, optimal, tree*, etc.
- Maximum feature size limited for efficiency

```
from sklearn.feature_extraction.text import TfidfVectorizer

tfidf = TfidfVectorizer(
    max_features=500,
    ngram_range=(1,2),
    stop_words='english',
    min_df=3,
    max_df=0.8,
    sublinear_tf=True
)
```

## 4.2 Handcrafted Features

To improve performance, several handcrafted features were added:

- Total text length
- Number of symbols (+ - \* / = < >)
- Number of lines
- Word count
- Sentence count
- Presence of algorithmic keywords
- Average word length

These features help models capture **structural complexity** beyond raw text. The final feature vector is a **concatenation of TF-IDF and handcrafted features**.

## 5. Classification Models

The goal of classification is to predict:

**Easy / Medium / Hard**

The following models were trained and evaluated:

## Models Tried

- Logistic Regression
- Linear Support Vector Machine (LinearSVC)
- Random Forest Classifier
- Histogram Gradient Boosting Classifier

## Best Model

- Random Forest Classifier
  - Class weighting used to handle imbalance
  - Achieved best accuracy on test set

## Evaluation Metrics

- Accuracy
- Precision, Recall, F1-Score
- Confusion Matrix

## 6. Regression Models

The regression task predicts a **continuous difficulty score**.

## Models Tried

- Ridge Regression
- Random Forest Regressor
- Gradient Boosting Regressor
- HistGradientBoostingRegressor

## Best Model

- GradientBoostingRegressor

Top 20 Most Important Features:	
1. number_count	: 0.0250
2. line_count	: 0.0224
3. word_count	: 0.0220
4. text_len	: 0.0177
5. avg_word_length	: 0.0140
6. possible	: 0.0134
7. sentence_count	: 0.0125
8. single	: 0.0115
9. input	: 0.0106
10. output	: 0.0086
11. integer	: 0.0081
12. symbol_count	: 0.0080
13. single_line	: 0.0075
14. 10	: 0.0073
15. line_input	: 0.0073
16. write	: 0.0072
17. contains	: 0.0069
18. integers	: 0.0069
19. given	: 0.0068
20. leq	: 0.0063

Training: Random Forest					
Accuracy: 0.5407 (54.07%)					
Classification Report:					
	precision	recall	f1-score	support	
easy	0.44	0.44	0.44	153	
hard	0.60	0.78	0.68	389	
medium	0.46	0.26	0.33	281	
accuracy			0.54	823	
macro avg	0.50	0.49	0.48	823	
weighted avg	0.52	0.54	0.51	823	
Confusion Matrix:					
[[ 68 48 37] [ 37 305 47] [ 50 159 72]]					

Chosen due to:

- Lower MAE and RMSE
- Better generalization
- Faster training on dense features

```
=====  
Training: Gradient Boosting  
=====  
  
Train MAE: 0.8707 | Test MAE: 1.6613  
Train RMSE: 1.0535 | Test RMSE: 2.0141  
Train R2: 0.7645 | Test R2: 0.1640
```

## Evaluation Metrics

- Mean Absolute Error (MAE)
- Root Mean Squared Error (RMSE)
- R<sup>2</sup> Score

## 7. Experimental Results

### 7.1 Classification Results

- Accuracy: ~54%
- Random Forest performed best among all classifiers
- Confusion matrix showed most confusion between Medium and Hard classes

### 7.2 Regression Results

- MAE ≈ 1.7
- RMSE ≈ 2.0
- R<sup>2</sup> ≈ 0.16

Although regression performance is moderate, the predicted scores align reasonably with difficulty trends.

## 8. Web Interface


A **Streamlit web application** was developed to demonstrate the system.

**Features:**

- User inputs problem text
- Displays:
  - Predicted difficulty score
  - Predicted difficulty class
- Loads trained models from disk
- Runs locally without internet access

## Technologies Used:

- Streamlit
- scikit-learn
- NumPy
- joblib



Screenshots of the web interface and sample predictions are included.

## 9. Conclusion

This project demonstrates that **problem difficulty can be predicted using textual information alone**. By combining NLP features with handcrafted complexity indicators, both classification and regression models were successfully trained.

AutoJudge provides a strong baseline for automated problem difficulty assessment.

